

Spring Boot Auto-configuration for Apache Geode & Pivotal GemFire

This lab is a modified version of the guide provided by John Blum at [apache-geode-docs](https://dzone.com/articles/spring-boot-auto-configuration-for-apache-geode)

This lab walks you through building a simple Customer Service, Spring Boot application using Apache Geode to manage Customer interactions. You should already be familiar with Spring Boot and Apache Geode.

By the end of this lesson, you should have a better understanding of what Spring Boot for Apache Geode's (SBDG) *auto-configuration* support actually does.

This lab compliments the [Auto-configuration vs. Annotation-based configuration](#) chapter with concrete examples.

This lab builds on the [Simplifying Apache Geode with Spring Data](#) presentation by John Blum during the 2017 SpringOne Platform conference. While this example as well as the example presented in the talk both use Spring Boot, only this example is using Spring Boot for Apache Geode (SBDG). This guide improves on the example from the presentation by using SBDG.

link:<https://docs.spring.io/spring-boot-data-geode-build/current/reference/html5/guides/boot-configuration.html>

link:<https://docs.spring.io/spring-boot-data-geode-build/current/reference/html5/index.html#geode-samples>

Application Domain Classes

We will build the Spring Boot, Customer Service application from the ground up.

Customer class

Like any sensible application development project, we begin by modeling the data our application needs to manage, namely a Customer. For this example, the Customer class is implemented as follows:

Customer class.

```
@Region("Customers")
@EqualsAndHashCode
@ToString(of = "name")
@RequiredArgsConstructor(staticName = "newCustomer")
public class Customer {

    @Id @NonNull @Getter
    private Long id;

    @NonNull @Getter
    private String name;
}
```

The Customer class uses [Project Lombok](#) to simplify the implementation so we can focus on the details we care about. Lombok is useful for testing or prototyping purposes. However, using Project Lombok is optional and in most production applications, and I would not recommend it.

Additionally, the Customer class is annotated with Spring Data Geode's (SDG) @Region annotation. @Region is a mapping annotation declaring the Apache Geode cache Region in which Customer data will be persisted.

Finally, the org.springframework.data.annotation.Id annotation was used to designate the Customer.id field as the identifier for Customer objects. The identifier is the Key used in the Entry stored in the "Customers" Region. A Region is a distributed version of java.util.Map.

If the @Region annotation is not explicitly declared, then SDG uses the simple name of the class, which in this case is "Customer", to identify the Region. However, there is another reason we explicitly annotated the Customer class with @Region, which we will cover below.

CustomerRepository interface

Next, we create a *Data Access Object* (DAO) to persist Customers to Apache Geode. We create the DAO using Spring Data's *Repository* abstraction:

CustomerRepository interface.

```
public interface CustomerRepository extends CrudRepository<Customer, Long> {

    Customer findByNameLike(String name);

}
```

CustomerRepository is a Spring Data *CrudRepository*. *CrudRepository* provides basic CRUD (CREATE, READ, UPDATE, and DELETE) data access operations along with the ability to define simple queries on Customers.

Spring Data Geode will create a proxy implementation for your application-specific *Repository* interfaces, implementing any query methods you may have explicitly defined on the interface in addition to the data access operations provided in the *CrudRepository* interface extension.

In addition to the base *CrudRepository* operations, *CustomerRepository* has additionally defined a *findByNameLike(:String):Customer* query method. The Apache Geode OQL query is derived from the method declaration.

Though it is beyond the scope of this document, Spring Data's *Repository* infrastructure is capable of generating data store specific queries (e.g. Apache Geode OQL) for *Repository* interface query method declarations just by introspecting the method signature. The query methods must conform to specific conventions. Alternatively, users may use *@Query* to annotate query methods to specify the raw query instead (i.e. OQL for Apache Geode, SQL for JDBC, possibly HQL for JPA, and so on).

CustomerServiceApplication (Spring Boot main class)

Now that we have created the basic domain classes of our Customer Service application, we need a main application class to drive the interactions with Customers:

CustomerServiceApplication class.

```
@SpringBootApplication
@EnableClusterAware
@EnableEntityDefinedRegions(basePackageClasses = Customer.class)
public class CustomerServiceApplication {

    public static void main(String[] args) {

        new SpringApplicationBuilder(CustomerServiceApplication.class)
            .web(WebApplicationType.NONE)
            .build()
            .run(args);
    }

    @Bean
    ApplicationRunner runner(CustomerRepository customerRepository) {

        return args -> {

            assertThat(customerRepository.count()).isEqualTo(0);

            Customer jonDoe = Customer.newCustomer(1L, "Jon Doe");

            System.err.printf("Saving Customer [%s]\n", jonDoe);

            jonDoe = customerRepository.save(jonDoe);

            assertThat(jonDoe).isNotNull();
            assertThat(jonDoe.getId()).isEqualTo(1L);
            assertThat(jonDoe.getName()).isEqualTo("Jon Doe");
            assertThat(customerRepository.count()).isEqualTo(1);

            System.err.println("Querying for Customer [SELECT * FROM /Customers WHERE name LIKE '%Doe']");

            Customer queriedJonDoe = customerRepository.findByNameLike("%Doe");

            assertThat(queriedJonDoe).isEqualTo(jonDoe);
        }
    }
}
```

```

        System.err.printf("Customer was [%s]%n", queriedJonDoe);
    };
}
}

```

The `CustomerServiceApplication` class is annotated with `@SpringBootApplication`. Therefore, the main class is a proper Spring Boot application equipped with all the features of Spring Boot (e.g. *auto-configuration*).

Additionally, we use Spring Boot's `SpringApplicationBuilder` in the main method to configure and bootstrap the Customer Service application.

Then, we declare a Spring Boot `ApplicationRunner` bean, which is invoked by Spring Boot after the Spring container (i.e. `ApplicationContext`) has been properly initialized and started. Our `ApplicationRunner` defines the Customer interactions performed by our Customer Service application.

Specifically, the runner creates a new `Customer` object ("Jon Doe"), saves him to the "Customers" Region, and then queries for "Jon Doe" using an OQL query with the predicate: `name LIKE '%Doe'`.

`%` is the wildcard for OQL text searches.

Running the Example

You can run the `CustomerServiceApplication` class from your IDE (e.g. Spring Tool Suite or IntelliJ IDEA) or from the command-line with the `mvn clean package` followed by the command.

There is nothing special you must do to run the `CustomerServiceApplication` class from inside your IDE. Simply create a run profile configuration and run it.

There is also nothing special about running the `CustomerServiceApplication` class from the command-line using `mvn`. Simply execute it with `spring-boot:bootRun`:

```
`$ mvn spring-boot:run
```

If you wish to adjust the log levels for either Apache Geode or Spring Boot while running the example, then you can set the log level for the individual Loggers (i.e. `org.apache` or `org.springframework`) in `src/main/resources/logback.xml`:

spring-geode-samples/boot/configuration/src/main/resources/logback.xml.

```
include:../Users/wlund/Dropbox/git-workspace/wxlund/spring-geode-workshop/configuration/src/main/resources/log
```

Auto-configuration for Apache Geode, Take One

"With great power comes great responsibility." - Uncle Ben

While it is not apparent (yet), there is a lot of hidden, intrinsic power provided by Spring Boot Data Geode (SBDG) in this example.

Cache instance

First, in order to put anything into Apache Geode you need a cache instance. A cache instance is also required to create `Regions` which ultimately store the application's data (state). Again, a `Region` is just a `Key/Value` data structure, like `java.util.Map`, mapping a `Key` to a `Value`, or an `Object`. A `Region` is actually much more than a simple `Map` since it is distributed. However, since `Region` implements `java.util.Map`, it can be treated as such.

A complete discussion of `Region` and its concepts are beyond the scope of this document. You may learn more by reading Apache Geode's User guide on [Developing with Apache Geode](#)

SBDG is opinionated and assumes most Apache Geode applications will be client applications in Apache Geode's [Topologies and Communication](#). Therefore, SBDG auto-configures a `ClientCache` instance by default.

The intrinsic `ClientCache auto-configuration` provided by SBDG can be made apparent by disabling it:

Disabling ClientCache Auto-configuration.

```

@SpringBootApplication(exclude = ClientCacheAutoConfiguration.class)
@EnableEntityDefinedRegions(basePackageClasses = Customer.class, clientRegionShortcut = ClientRegionShortcut

```

```
public class CustomerServiceApplication {
    // ...
}
```

Note the `exclude` on the `ClientCacheAutoConfiguration.class`.

With the correct log level set, you will see an error message similar to:

Error resulting from no `ClientCache` instance.

2020-08-24 17:26:20,964 ERROR agnostics.LoggingFailureAnalysisReporter: 40 -

APPLICATION FAILED TO START

Description:

Parameter 0 of method runner in `com.vmware.spring.geode.workshop.crm.CustomerServiceApplication` required a bean of type '`com.vmware.spring.geode.workshop.crm.repo.CustomerRepository`' that could not be found.

Action:

Consider defining a bean of type '`com.vmware.spring.geode.workshop.crm.repo.CustomerRepository`' in your configuration.

Essentially, the `CustomerRepository` could not be injected into our `CustomerServiceApplication` class, `ApplicationRunner` bean method because the `CustomerRepository`, which depends on the "Customers" Region, could not be created. The `CustomerRepository` could not be created because the "Customers" Region could not be created. The "Customers" Region could not be created because there was no cache instance available (e.g. `ClientCache`) to create Regions, resulting in a trickling effect.

We are going to take advantage of an annotation that allows our developer environment to determine what environment we are running in and create a local (near) cache or a client server connection through auto configuration via the following:

Equivalent `ClientCache` configuration.

```
@SpringBootApplication
@EnableClusterAware
@EnableEntityDefinedRegions(basePackageClasses = Customer.class)
public class CustomerServiceApplication {
    // ...
}
```

Otherwise you would need to explicitly declare the `@ClientCacheApplication` annotation if you were not using SBDG.

Repository instance

We are also using the Spring Data (Geode) *Repository* infrastructure in the Customer Service application. This should be evident from our declaration and definition of the application-specific `CustomerRepository` interface.

If we disable the Spring Data *Repository* auto-configuration:

Disabling Spring Data Repositories Auto-configuration.

```
@SpringBootApplication(exclude = RepositoriesAutoConfiguration.class)
@EnableEntityDefinedRegions(basePackageClasses = Customer.class, clientRegionShortcut = ClientRegionShortcut)
public class CustomerServiceApplication {
    // ...
}
```

The application would throw a similar error on startup:

2020-08-24 17:38:49,107 INFO xt.annotation.ConfigurationClassEnhancer: 323 - @Bean method `PdxConfiguration.pdxDiskStoreAwareBeanFactoryPostProcessor` is non-static and returns an object assignable to Spring's `BeanFactoryPostProcessor` interface. This will result in a failure to process annotations such as `@Autowired`, `@Resource` and `@PostConstruct` within the method's declaring `@Configuration` class. Add the 'static' modifier to this method to avoid these container lifecycle issues; see `@Bean` javadoc for complete details. 2020-08-24 17:38:49,112 INFO xt.annotation.ConfigurationClassEnhancer: 323 - @Bean method `RegionTemplateAutoConfiguration.regionTemplateBeanFactoryPostProcessor` is non-static and returns an object assignable to Spring's `BeanFactoryPostProcessor` interface. This will result in a failure to process annotations such as `@Autowired`,

@Resource and @PostConstruct within the method's declaring @Configuration class. Add the 'static' modifier to this method to avoid these container lifecycle issues; see @Bean javadoc for complete details. 2020-08-24 17:38:49,184 INFO

trationDelegate\$BeanPostProcessorChecker: 335 - Bean

'org.springframework.geode.boot.autoconfigure.DataImportExportAutoConfiguration' of type

[org.springframework.geode.boot.autoconfigure.DataImportExportAutoConfiguration\$\$EnhancerBySpringCGLIB\$\$cb0836d1] is not eligible for getting processed by all BeanPostProcessors (for example: not eligible for auto-proxying) 2020-08-24 17:38:49,229 WARN

ation.AnnotationConfigApplicationContext: 559 - Exception encountered during context initialization - cancelling refresh attempt:

org.springframework.beans.factory.UnsatisfiedDependencyException: Error creating bean with name 'runner' defined in

com.vmware.spring.geode.workshop.crm.CustomerServiceApplication: Unsatisfied dependency expressed through method 'runner'

parameter 0; nested exception is org.springframework.beans.factory.NoSuchBeanDefinitionException: No qualifying bean of type

'com.vmware.spring.geode.workshop.crm.repo.CustomerRepository' available: expected at least 1 bean which qualifies as autowire

candidate. Dependency annotations: {} 2020-08-24 17:38:49,236 INFO ConditionEvaluationReportLoggingListener: 136 -

Error starting ApplicationContext. To display the conditions report re-run your application with 'debug' enabled. 2020-08-24

17:38:49,307 ERROR agnostics.LoggingFailureAnalysisReporter: 40 -

APPLICATION FAILED TO START

Description:

Parameter 0 of method runner in com.vmware.spring.geode.workshop.crm.CustomerServiceApplication required a bean of type 'com.vmware.spring.geode.workshop.crm.repo.CustomerRepository' that could not be found.

Action:

Consider defining a bean of type 'com.vmware.spring.geode.workshop.crm.repo.CustomerRepository' in your configuration.

In this case, there was simply no proxy implementation for the `CustomerRepository` interface provided by the framework since the *auto-configuration* was disabled. The `ClientCache` and "Customers" Region do exist in this case, though.

The Spring Data *Repository auto-configuration* even takes care of locating our application *Repository* interface definitions for us.

Without *auto-configuration*, you would need to explicitly:

Equivalent Spring Data Repositories configuration.

```
@SpringBootApplication(exclude = RepositoriesAutoConfiguration.class)
@EnableEntityDefinedRegions(basePackageClasses = Customer.class, clientRegionShortcut = ClientRegionShortcut)
@EnableGemfireRepositories(basePackageClasses = CustomerRepository.class)
public class CustomerServiceApplication {
    // ...
}
```

That is, you would need to explicitly declare the `@EnableGemfireRepositories` annotation and set the `basePackages` attribute, or the equivalent, type-safe `basePackageClasses` attribute, to the package containing your application *Repository* interfaces, if you were not using SBDG.

Entity-defined Regions

So far, the only explicit declaration of configuration in our Customer Service application is the `@EnableEntityDefinedRegions` annotation.

As was alluded to above, there was another reason we explicitly declared the `@Region` annotation on our `Customer` class.

Using SDG's `@EnableEntityDefinedRegions` annotation is very convenient and can scan for the Regions (whether client or server (peer) Regions) required by your application based the entity classes themselves (e.g. `Customer`). In addition, we are going to take advantage of a very useful annotation for allowing SBDG to automatically detect our topology (e.g. local/near cache or client server):

Annotation-based config for the "Customers" Region.

```
@SpringBootApplication
@EnableClusterAware
@EnableEntityDefinedRegions(basePackageClasses = Customer.class)
class CustomerServiceApplication { }
```

The `basePackageClasses` attribute is an alternative to `basePackages`, and a type-safe way to target the packages (and subpackages) containing the entity classes that your application will persist to Apache Geode. You only need to choose one class from each top-level package for where you want the scan to begin. Spring Data Geode uses this class to determine the package to begin the scan. `'basePackageClasses'` accepts an array of `Class` types so you can specify multiple independent top-level packages. The annotation also includes the ability to filter types.

However, the `@EnableEntityDefinedRegions` annotation only works when the entity class (e.g. `Customer`) is explicitly annotated with the `@Region` annotation (e.g. `@Region("Customers")`), otherwise it ignores the class. You will find this annotation on `Customer` in the model package.

Initially we just want to get up and running as quickly as possible, without a lot of ceremony and fuss. By not starting a locator/region and using `@EnableClusterAware` we allow SBDG to detect that we are using a client `LOCAL` Region and we are not required to start a cluster of servers for the client to be able to store data. This allows rapid development for developers.

While client `LOCAL` Regions can be useful for some purposes (e.g. local processing, querying and aggregating of data), it is more common for a client to persist data in a cluster of servers, and for that data to be shared by multiple clients (instances) in the application architecture, especially as the application is scaled out to handle demand.

Switching to Client/Server

We continue with our example by switching from a local context to a client/server topology.

If you are rapidly prototyping and developing your application and simply want to lift off the ground quickly, then it is useful to start locally and gradually migrate towards a client/server architecture.

To switch to client/server, all you need to do is startup a `gfsh` locator(s) and server(s) and SBDG will automatically detect your topology. `@EnableEntityDefinedRegions` annotation declaration:

Client/Server Topology Region Configuration.

```
@SpringBootApplication
@EnableClusterAware
@EnableEntityDefinedRegions(basePackageClasses = Customer.class)

class CustomerServiceApplication { }
```

The `ClusterAware` auto configuration determines the topology by determining whether there is a cluster of servers to communicate with and to store/access data from. Clearly, there are no servers or cluster running yet.

There are several ways in which to start a cluster. For example, you may use Spring to configure and bootstrap the cluster, which has been demonstrated [here](#).


Although, for this example, we are going to use the tools provided with Apache Geode, or VMWare Tanzu GemFire, i.e. *Gfsh* (`GemFire/Geode Shell`) for reasons that will become apparent later.

You need to [download](#) and [install](#) a full distribution of Apache Geode to make use of the provided tools. After installation, you will need to set the `GEODE` (or `GEMFIRE`) environment variable to the location of your installation. Additionally, add `$GEODE/bin` to your system `$PATH`.

Once Apache Geode has been successfully installed, you can open a command prompt (terminal) and do:

Running Gfsh.

```
$ gfsh
```



```
1.2.1
```

```
Monitor and Manage Apache Geode
gfsh>
```

You are set to go.

The lab environment running on kubernetes was discussed in the lab setup. Two locators and 2 servers have been started.

Execute the *Gfsh* shell script using:

With our gemfire-cluster create with an two Apache Geode Locators and (Cache) Servers running we can verify by listing and describing the members:

List and Describe Members.

```
gfsh>list members
      Name      | Id
-----|-----
LocatorOne     | 10.99.199.24(LocatorOne:68425:locator)<ec><v0>:1024
ServerOne      | 10.99.199.24(ServerOne:68434)<v1>:1025
```

```
gfsh>describe member --name=ServerOne
Name      : ServerOne
Id        : 10.99.199.24(ServerOne:68434)<v1>:1025
Host      : 10.99.199.24
Regions   :
PID       : 68434
Groups    :
Used Heap : 27M
Max Heap  : 3641M
Working Dir : /Users/jblum/pivdev/lab/ServerOne
Log file   : /Users/jblum/pivdev/lab/ServerOne/ServerOne.log
Locators   : 10.99.199.24[10334]
```

```
Cache Server Information
Server Bind      : null
Server Port     : 40404
Running         : true
Client Connections : 0
```

What happens if we try to run the application now?

RegionNotFoundException.

```
17:42:16.873 [main] ERROR o.s.b.SpringApplication - Application run failed
java.lang.IllegalStateException: Failed to execute ApplicationRunner
...
    at example.app.crm.CustomerServiceApplication.main(CustomerServiceApplication.java:51) [classes/?:?]
Caused by: org.springframework.dao.DataAccessResourceFailureException: remote server on 10.99.199.24(SpringBasedCacheClient)
    at org.springframework.data.gemfire.GemfireCacheUtils.convertGemfireAccessException(GemfireCacheUtils.java:90)
    at org.springframework.data.gemfire.GemfireAccessor.convertGemfireAccessException(GemfireAccessor.java:90)
    at org.springframework.data.gemfire.GemfireTemplate.find(GemfireTemplate.java:329) ~[spring-data-gemfire-2.0.9.jar:2.0.9]
    at org.springframework.data.gemfire.repository.support.SimpleGemfireRepository.count(SimpleGemfireRepository.java:100)
    ...
    at example.app.crm.CustomerServiceApplication.lambda$runner$0(CustomerServiceApplication.java:59) ~[classes/?:?]
    at org.springframework.boot.SpringApplication.callRunner(SpringApplication.java:783) ~[spring-boot-2.0.9.jar:2.0.9]
    ... 3 more
Caused by: org.apache.geode.cache.client.ServerOperationException: remote server on 10.99.199.24(SpringBasedCacheClient)
    at org.apache.geode.cache.client.internal.AbstractOp.processChunkedResponse(AbstractOp.java:352) ~[geode-core-1.2.1.jar:1.2.1]
    at org.apache.geode.cache.client.internal.QueryOp$QueryOpImpl.processResponse(QueryOp.java:170) ~[geode-core-1.2.1.jar:1.2.1]
    at org.apache.geode.cache.client.internal.AbstractOp.processResponse(AbstractOp.java:230) ~[geode-core-1.2.1.jar:1.2.1]
    at org.apache.geode.cache.client.internal.AbstractOp.attempt(AbstractOp.java:394) ~[geode-core-1.2.1.jar:1.2.1]
    at org.apache.geode.cache.client.internal.AbstractOp.attemptReadResponse(AbstractOp.java:203) ~[geode-core-1.2.1.jar:1.2.1]
    at org.apache.geode.cache.client.internal.ConnectionImpl.execute(ConnectionImpl.java:275) ~[geode-core-1.2.1.jar:1.2.1]
    at org.apache.geode.cache.client.internal.pooling.PooledConnection.execute(PooledConnection.java:332) ~[geode-core-1.2.1.jar:1.2.1]
    at org.apache.geode.cache.client.internal.OpExecutorImpl.executeWithPossibleReAuthentication(OpExecutorImpl.java:158) ~[geode-core-1.2.1.jar:1.2.1]
    at org.apache.geode.cache.client.internal.OpExecutorImpl.execute(OpExecutorImpl.java:115) ~[geode-core-1.2.1.jar:1.2.1]
    at org.apache.geode.cache.client.internal.PoolImpl.execute(PoolImpl.java:763) ~[geode-core-1.2.1.jar:1.2.1]
    at org.apache.geode.cache.client.internal.QueryOp.execute(QueryOp.java:58) ~[geode-core-1.2.1.jar:1.2.1]
    at org.apache.geode.cache.client.internal.ServerProxy.query(ServerProxy.java:70) ~[geode-core-1.2.1.jar:1.2.1]
    at org.apache.geode.cache.query.internal.DefaultQuery.executeOnServer(DefaultQuery.java:456) ~[geode-core-1.2.1.jar:1.2.1]
    at org.apache.geode.cache.query.internal.DefaultQuery.execute(DefaultQuery.java:338) ~[geode-core-1.2.1.jar:1.2.1]
    at org.springframework.data.gemfire.GemfireTemplate.find(GemfireTemplate.java:311) ~[spring-data-gemfire-2.0.9.jar:2.0.9]
    at org.springframework.data.gemfire.repository.support.SimpleGemfireRepository.count(SimpleGemfireRepository.java:100)
    ...
    at example.app.crm.CustomerServiceApplication.lambda$runner$0(CustomerServiceApplication.java:59) ~[classes/?:?]
    at org.springframework.boot.SpringApplication.callRunner(SpringApplication.java:783) ~[spring-boot-2.0.9.jar:2.0.9]
    ... 3 more
```

```
Caused by: org.apache.geode.cache.query.RegionNotFoundException: Region not found: /Customers
    at org.apache.geode.cache.query.internal.DefaultQuery.checkQueryOnPR(DefaultQuery.java:599) ~[geode-core-1.2.1.jar:1.2.1]
    at org.apache.geode.cache.query.internal.DefaultQuery.execute(DefaultQuery.java:348) ~[geode-core-1.2.1.jar:1.2.1]
    at org.apache.geode.cache.query.internal.DefaultQuery.execute(DefaultQuery.java:319) ~[geode-core-1.2.1.jar:1.2.1]
    at org.apache.geode.internal.cache.tier.sockets.BaseCommandQuery.processQueryUsingParams(BaseCommandQuery.java:104) ~[geode-core-1.2.1.jar:1.2.1]
    at org.apache.geode.internal.cache.tier.sockets.BaseCommandQuery.processQuery(BaseCommandQuery.java:65) ~[geode-core-1.2.1.jar:1.2.1]
    at org.apache.geode.internal.cache.tier.sockets.command.Query.cmdExecute(Query.java:91) ~[geode-core-1.2.1.jar:1.2.1]
    at org.apache.geode.internal.cache.tier.sockets.BaseCommand.execute(BaseCommand.java:165) ~[geode-core-1.2.1.jar:1.2.1]
    at org.apache.geode.internal.cache.tier.sockets.ServerConnection.doNormalMsg(ServerConnection.java:791) ~[geode-core-1.2.1.jar:1.2.1]
    at org.apache.geode.internal.cache.tier.sockets.ServerConnection.doOneMessage(ServerConnection.java:922) ~[geode-core-1.2.1.jar:1.2.1]
    at org.apache.geode.internal.cache.tier.sockets.ServerConnection.run(ServerConnection.java:1180) ~[geode-core-1.2.1.jar:1.2.1]
    ...

```

The application fails to run because we (deliberately) did not create a corresponding, server-side, "Customers" Region. In order for a client to send data via a client `PROXY` Region (a Region with no local state) to a server in a cluster, at least one server in the cluster must have a matching Region by name (i.e. "Customers").

Indeed, there are no Regions in the cluster:

List Regions.

```
gfsh>list regions
No Regions Found
```

Of course, you could create the matching server-side, "Customers" Region using *Gfsh*:

```
gfs# create region --name=Customers --type=PARTITION
```

But, what if you have hundreds of application domain objects each requiring a Region for persistence? It is not an unusual or unreasonable requirement in any practical enterprise scale application.

While it is not a "convention" in Spring Boot for Apache Geode (SBDG), Spring Data for Apache Geode (SDG) comes to our rescue. We simply only need to enable cluster configuration from the client:

Enable Cluster Configuration.

```
@SpringBootApplication
@EnableClusterAware
@EnableEntityDefinedRegions(basePackageClasses = Customer.class)    @SpringBootApplication
public class CustomerServiceApplication {
    // ...
}
```

SDG is careful not to stomp on existing Regions since those Regions may have data already. Declaring the `@EnableClusterConfiguration`` annotation is a useful development-time feature, but it is recommended that you explicitly define and declare your Regions in production environments, either using *Gfsh* or Spring *config*.

It is now possible to replace the SDG `@EnableClusterConfiguration` annotation with SBDG's `@EnableClusterAware` annotation as you see above in the src code, which has the same effect of pushing configuration metadata from the client to the server (or cluster). Additionally, SBDG's `@EnableClusterAware` annotation makes it unnecessary to explicitly have to configure parameters like `clientRegionShortcut` on the SDG `@EnableEntityDefinedRegions` annotation (or similar annotation, e.g. SDG's `@EnableCachingDefinedRegions`). Finally, because the SBDG's `@EnableClusterAware` annotation is meta-annotated with SDG's `@EnableClusterConfiguration` annotation is automatically configured with the `useHttp` attribute to `true`.

Now, we can run our application again, and this time, it works!

Client/Server Run Successful.

```

      .
  /\  /_____,
 (  ) \_____,
  \  /_____,
   |_____,
=====|_|=====|_|/=/=/=/=/
:: Spring Boot ::      (v2.0.9.RELEASE)

```

```
Saving Customer [Customer(name=Jon Doe)]
Querying for Customer [SELECT * FROM /Customers WHERE name LIKE '%Doe']
Customer was [Customer(name=Jon Doe)]
```


Process finished with exit code 0

In the cluster (server-side), we will also see that the "Customers" Region was created successfully:

List & Describe Regions.

```
gfsh>list regions
List of regions
-----
Customers

gfsh>describe region --name=/Customers
.....
Name           : Customers
Data Policy    : partition
Hosting Members : ServerOne
```

Non-Default Attributes Shared By Hosting Members

Type	Name	Value
Region	size	1
	data-policy	PARTITION

We see that the "Customers" Region has a size of 1, containing "Jon Doe".

We can verify this by querying the "Customers" Region:

Query for all Customers.

```
gfsh>query --query="SELECT customer.name FROM /Customers customer"
Result : true
Limit  : 100
Rows   : 1

Result
-----
Jon Doe
```

That was easy!

Auto-configuration for Apache Geode, Take Two

What may not be apparent in this example up to this point is how the data got from the client to the server. Certainly, our client did send Jon Doe to the server, but our Customer class is not `java.io.Serializable`. So, how was an instance of Customer streamed and sent from the client to the server then (it is using a Socket)?

Any object sent over a network, between two Java processes, or streamed to/from disk, must be serializable, no exceptions!

Furthermore, when we started the cluster, we did not include any application domain classes on the classpath of any server in the cluster.

As further evidence, we can adjust our query slightly:

Invalid Query.

```
gfsh>query --query="SELECT * FROM /Customers"
Message : Could not create an instance of a class example.app.crm.model.Customer
Result  : false
```

If you tried to perform a get, you would hit a similar error:

Region.get(key).

```
gfsh>get --region=/Customers --key=1 --key-class=java.lang.Long
Message : Could not create an instance of a class example.app.crm.model.Customer
Result  : false
```

So, how was the data sent then? How were we able to access the data stored in the server(s) on the cluster with the OQL query `SELECT customer.name FROM /Customers customer` as seen above?

Well, Apache Geode and Pivotal GemFire provide 2 proprietary serialization formats in addition to *Java Serialization*: {apache-geode-docs}/developing/data_serialization/gemfire_data_serialization.html[Data Serialization] and {apache-geode-docs}/developing/data_serialization/gemfire_pdx_serialization.html[PDX], or *Portable Data Exchange*.

While *Data Serialization* is more efficient, PDX is more flexible (i.e. "portable"). PDX enables data to be queried in serialized form and is the format used to support both Java and Native Clients (C++, C#) simultaneously. Therefore, PDX is auto-configured in Spring Boot Data Geode (SBDG) by default.

This is convenient since you may not want to implement `java.io.Serializable` for all your application domain model types that you store in Apache Geode. In other cases, you may not even have control over the types referred to by your application domain model types to make them `Serializable`, such as when using a 3rd party library.

So, SBDG auto-configures PDX and uses Spring Data Geode's `MappingPdxSerializer` as the `PdxSerializer` to de/serialize all application domain model types.

If we were to disable PDX *auto-configuration*, you would see the effects of trying to serialize a non-serializable type, `Customer`. This would be apparent when you tried to query `/Customers`. We will leave this as an exercise for the student.

SBDG takes care of all your serialization needs without you having to configure serialization or implement `java.io.Serializable` in all your application domain model types, including types your application domain model types might refer to, which may not be possible.

If you were not using SBDG, then you would need to enable PDX serialization explicitly.

The Spring Boot Data Geode/Gemfire PDX *auto-configuration* provided by SBDG is equivalent to:

Equivalent PDX Configuration with Spring Data Gemfire.

```
@SpringBootApplication
@ClientCacheApplication
@EnableEntityDefinedRegions(basePackageClasses = Customer.class)
@EnableClusterConfiguration(useHttp = true)
@EnablePdx
public class CustomerServiceApplication {
    // ...
}
```

In addition to the `@ClientCacheApplication` annotation, you would need to annotate the `CustomerServiceApplication` class with SDG's `@EnablePdx` annotation, which is responsible for configuring PDX serialization and registering SDG's `MappingPdxSerializer`.

[Sample Applications.](#)